

Scientific Programming

Practical 8

Introduction

Luca Bianco - Academic Year 2019-20
luca.bianco@fmach.it

Pandas

Pandas (**panel-data**) is a very efficient library to deal with **numerical tables** and time series

Two data structures:

Series: 1D tables

DataFrames: 2D tables

<https://pandas.pydata.org/>

Series

Series are 1-dimensional structures (like lists) containing data. Series are characterized by two types of information: the **values** and the **index** (a list of labels associated to the data). A bit like **list** and a bit like **dictionary**!

```
Values and index explicitly defined
A    15
B     7
C    20
D     3
E    15
F     1
G     5
H    17
I    15
L    17
dtype: int64
The index: Index(['A', 'B', 'C', 'D', 'E', 'F', 'G', 'H', 'I', 'L'], dtype='object')
The values: [15  7 20  3 15  1  5 17 15 17]
-----

From dictionary
forty    40
four     4
one      1
ten     10
three    3
two      2
dtype: int64
Index(['forty', 'four', 'one', 'ten', 'three', 'two'], dtype='object')
[40  4  1 10  3  2]
```

```
import pandas as pd
import random

print("Values and index explicitly defined")
#values and index explicitly defined
S = pd.Series([random.randint(0,20) for x in range(0,10)],
              index = list("ABCDEFGHIL"))

print(S)
print("The index:", S.index)
print("The values:", S.values)

print("-----\n")
print("From dictionary")
#from a dictionary
S1 = pd.Series({"one" : 1, "two" : 2, "ten": 10,
               "three" : 3, "four": 4, "forty" : 40})

print(S1)
print(S1.index)
print(S1.values)
```

Series

Series are 1-dimensional structures (like lists) containing data. Series are characterized by two types of information: the **values** and the **index** (a list of labels associated to the data). A bit like **list** and a bit like **dictionary**!

Default index

```
0    8
1    2
2    8
3   10
4    1
5    5
6    3
7    8
8    9
9    5
```

dtype: int64

RangeIndex(start=0, stop=10, step=1)

```
[ 8  2  8 10  1  5  3  8  9  5]
```

Same value repeated

```
0    1.27
1    1.27
2    1.27
3    1.27
4    1.27
5    1.27
6    1.27
7    1.27
8    1.27
9    1.27
```

dtype: float64

RangeIndex(start=0, stop=10, step=1)

```
[1.27 1.27 1.27 1.27 1.27 1.27 1.27 1.27 1.27 1.27]
```

```
print("Default index")
#index added by default
myData = [random.randint(0,10) for x in range(10)]
S2 = pd.Series(myData)
```

```
print(S2)
print(S2.index)
print(S2.values)
```

```
print("-----\n")
print("Same value repeated")
S3 = pd.Series(1.27, range(10))
print(S3)
print(S3.index)
print(S3.values)
```

Series

Data in a series can be accessed by using the **label** (i.e. the index) as in a dictionary or through its **position** as in a list. Slicing is also allowed both by **position** and **index**.

In the latter case, `Series[S:E]` with **S** and **E** **indexes**, both **S** and **E** **are included**.

```
import pandas as pd
import random

#values and index explicitly defined
S = pd.Series([random.randint(0,20) for x in range(0,10)],
              index = list("ABCDEFGHIL"))

print(S)
print("")

print("Value at label \"A\":", S["A"])
print("Value at index 1:", S[1])
print("")

print("Slicing from 1 to 3:") #note 3 excluded
print(S[1:3])
print("")
print("Slicing from C to H:") #note H included!
print(S["C":"H"])
print("")

print("Retrieving from list:")
print(S[[1,3,5,7,9]])
print(S[["A","C","E","G"]])
print("")

print("Top 3")
print(S.head(3))
print("")
print("Bottom 3")
print(S.tail(3))
```

```
A    15
B    11
C     4
D     7
E     1
F     4
G    15
H    14
I    14
L    17
dtype: int64
```

```
Value at label "A": 15
Value at index 1: 11
```

```
Slicing from 1 to 3:
B    11
C     4
dtype: int64
```

```
Slicing from C to H:
C     4
D     7
E     1
F     4
G    15
H    14
dtype: int64
```

Series

Data in a series can be accessed by using the **label** (i.e. the index) as in a dictionary or through its **position** as in a list. Slicing is also allowed both by **position** and **index**.

In the latter case, `Series[S:E]` with **S** and **E** **labels**, both **S** and **E** are included.

```
import pandas as pd
import random

#values and index explicitly defined
S = pd.Series([random.randint(0,20) for x in range(0,10)],
              index = list("ABCDEFGHIL"))

print(S)
print("")

print("Value at label \"A\":", S["A"])
print("Value at index 1:", S[1])
print("")

print("Slicing from 1 to 3:") #note 3 excluded
print(S[1:3])
print("")
print("Slicing from C to H:") #note H included!
print(S["C":"H"])
print("")

print("Retrieving from list:")
print(S[[1,3,5,7,9]])
print(S[["A","C","E","G"]])
print("")

print("Top 3")
print(S.head(3))
print("")
print("Bottom 3")
print(S.tail(3))
```

Retrieving from list:

```
B    11
D     7
F     4
H    14
L    17
dtype: int64
A    15
C     4
E     1
G    15
dtype: int64
```

Top 3

```
A    15
B    11
C     4
dtype: int64
```

Bottom 3

```
H    14
I    14
L    17
dtype: int64
```

Series

Example: Given a list of 10 integers and we want to divide them by 2.

Important operations on series:

Operator broadcasting

Operations can automatically be broadcast to the entire Series. This is a quite cool feature and **saves us from looping through the elements of the Series.**

Without pandas:

```
import random

listS = [random.randint(0,20) for x in range(0,10)]

print(listS)

for el in range(0,len(listS)):
    listS[el] /=2  #compact of X = X / 2

print(listS)
```

[6, 4, 5, 19, 14, 16, 9, 3, 13, 11]
[3.0, 2.0, 2.5, 9.5, 7.0, 8.0, 4.5, 1.5, 6.5, 5.5]

Series

Example: Given a list of 10 integers and we want to divide them by 2.

Important operations on series:

Operator broadcasting

Operations can automatically be broadcast to the entire Series. This is a quite cool feature and **saves us from looping through the elements of the Series.**

With pandas (operator broadcasting):

```
A    4
B   13
C   14
D    6
E   13
F    2
G   13
H   19
I   20
L    7
```

dtype: int64

```
import pandas as pd
import random

S = pd.Series([random.randint(0,20) for x in range(0,10)],
               index = list("ABCDEFGHIL"))

print(S)
print("")
S1 = S / 2
print(S1)
```

```
A    2.0
B    6.5
C    7.0
D    3.0
E    6.5
F    1.0
G    6.5
H    9.5
I   10.0
L    3.5
```

dtype: float64

Series

Important operations on series:

Operator broadcasting

Filtering

We can also apply boolean operators to obtain only the **sub-Series** with all the values satisfying a specific condition. This allows us to **filter** the Series.

```
import pandas as pd
import random

S = pd.Series([random.randint(0,20) for x in range(0,10)],
              index = list("ABCDEFGHIL"))

print(S)
print("")
S1 = S>10
print(S1)
print("")
S2 = S[S > 10]
print(S2)
```

```
A    3
B    3
C   18
D    1
E   12
F   11
G    4
H   11
I    5
L   14
dtype: int64
```

```
A    False
B    False
C     True
D    False
E     True
F     True
G    False
H     True
I    False
L     True
dtype: bool
```

```
C    18
E    12
F    11
H    11
L    14
dtype: int64
```



series of True and False
where condition is/is not
met

Series

Important operations on series:

Operator broadcasting

Filtering

Computing stats

```
import pandas as pd
import random

S = pd.Series([random.randint(0,10) for x in range(0,10)],
               index = list("ABCDEFGHIL"))

print("The data:")
print(S)
print("")
print("Its description")
print(S.describe())
print("")
print("Specifying different quantiles:")
print(S.quantile([0.1,0.2,0.8,0.9]))
print("")
print("Histogram:")
print(S.value_counts())
print("")
print("The type is a Series:")
print(type(S.value_counts()))
print("Summing the values:")
print(S.sum())
print("")
print("The cumulative sum:")
print(S.cumsum())
```

see notes for the complete
results and other features
like `Series.fillna(values)`

The data:

A	3
B	10
C	3
D	4
E	9
F	9
G	5
H	0
I	10
L	5

dtype: int64

Its description

count	10.000000
mean	5.800000
std	3.489667
min	0.000000
25%	3.250000
50%	5.000000
75%	9.000000
max	10.000000

dtype: float64

Specifying different quantiles:

0.1	2.7
0.2	3.0
0.8	9.2
0.9	10.0

dtype: float64

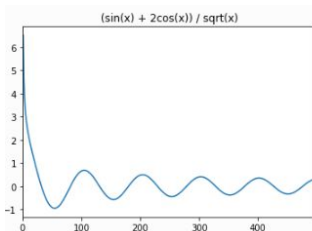
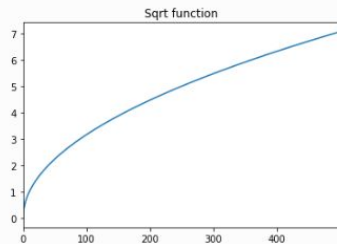
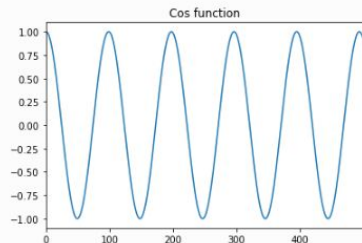
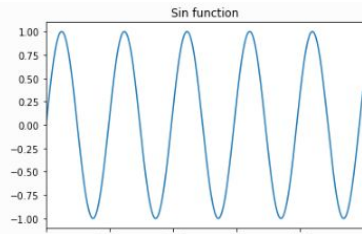
Histogram:

10	2
9	2
5	2
3	2
4	1
0	1

dtype: int64

Plotting data

It is quite easy to plot data in Series and DataFrames thanks to matplotlib



```
import math
import matplotlib.pyplot as plt
import pandas as pd
```

```
x = [i/10 for i in range(0,500)]
```

```
y = [math.sin(2*i/3.14 ) for i in x]
y1 = [math.cos(2*i/3.14 ) for i in x]
y2 = [math.sqrt(i) for i in x]
#print(x)
```

```
ySeries = pd.Series(y)
ySeries1 = pd.Series(y1)
ySeries2 = pd.Series(y2)
```

```
ySeries.plot()
plt.title("Sin function")
plt.show()
plt.close()
ySeries1.plot()
plt.title("Cos function")
plt.show()
plt.close()
ySeries2.plot()
plt.title("Sqrt function")
plt.show()
plt.close()
ySeries2 = (ySeries + 2*ySeries1)/ySeries2
ySeries2.plot()
plt.title("(sin(x) + 2cos(x)) / sqrt(x)")
plt.show()
```

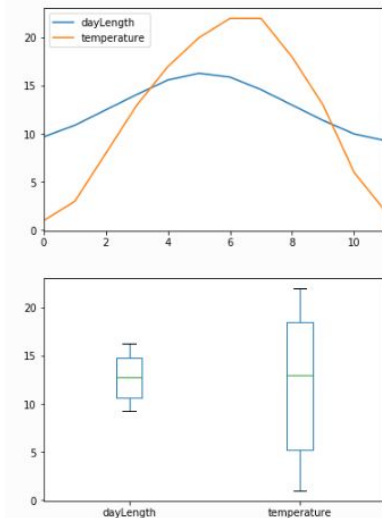
DataFrames

2D analogous of Series.

They have an **index** and several **columns**.

Data can be dishomogeneous.

Most of the the things seen for Series apply to DataFrames



```
import pandas as pd

myData = {
    "temperature" : pd.Series([1, 3, 8, 13, 17, 20, 22, 22, 18, 13, 6, 2],
                             index = ["Jan", "Feb", "Mar", "Apr", "May", "Jun",
                                       "Jul", "Aug", "Sep", "Oct", "Nov", "Dec"]),
    "dayLength" : pd.Series([9.7, 10.9, 12.5, 14.1, 15.6, 16.3, 15.9,
                            14.6, 13, 11.4, 10, 9.3],
                           index = ["Jan", "Feb", "Mar", "Apr", "May", "Jun",
                                    "Jul", "Aug", "Sep", "Oct", "Nov", "Dec"])
}

DF = pd.DataFrame(myData)
print(DF)

print(DF.columns)
print(DF.index)
```

	dayLength	temperature
Jan	9.7	1
Feb	10.9	3
Mar	12.5	8
Apr	14.1	13
May	15.6	17
Jun	16.3	20
Jul	15.9	22
Aug	14.6	22
Sep	13.0	18
Oct	11.4	13
Nov	10.0	6
Dec	9.3	2

```
Index(['dayLength', 'temperature'], dtype='object')
Index(['Jan', 'Feb', 'Mar', 'Apr', 'May', 'Jun', 'Jul', 'Aug', 'Sep', 'Oct',
       'Nov', 'Dec'],
      dtype='object')
```

DataFrames

We can **load external files**, extract info and apply operators, broadcasting and filtering...

Load from file



```
import pandas as pd
```

```
orders = pd.read_csv("file_samples/sampledata_orders.csv", sep=";",  
index_col =0, header=0)
```

```
print("The Order Quantity column (top 5)")
```

```
print(orders["Order Quantity"].head(5))
```

```
print("")
```

```
print("The Sales column (top 10)")
```

```
print(orders.Sales.head(10))
```

```
print("")
```

```
print("The row with ID:50")
```

```
r50 = orders.loc[50]
```

```
print(r50)
```

```
print("")
```

```
print("The third row:")
```

```
print(orders.iloc[3])
```

```
print("The Order Quantity, Sales, Discount and Profit of the 2nd,  
4th, 6th and 8th row:")
```

```
print(orders[1:8:2][["Order Quantity", "Sales", "Discount", "Profit"]])
```

```
print("The Order Quantity, Sales, Discount and Profit of orders with  
discount > 10%:")
```

```
print(orders[orders["Discount"] > 0.1][["Order Quantity", "Sales",  
"Discount", "Profit"]])
```

1. Select by column `DataFrame[col]` returns a Series
2. Select by row label `DataFrame.loc[row_label]` returns a Series
3. Select row by integer location `DataFrame.iloc[row_position]` returns a Series
4. Slice rows `DataFrame[S:E]` (S and E are labels, both included) returns a DataFrame
5. Select rows by boolean vector `DataFrame[bool_vect]` returns a DataFrame

Row ID	Sales	Profit	Product	Category
1	261.5400	-213.25	Office Supplies	
49	10123.0200	457.81	Office Supplies	
50	244.5700	46.71	Office Supplies	
80	4965.7595	1198.97	Technology	
85	394.2700	30.94	Office Supplies	
86	146.6900	4.43	Furniture	
97	93.5400	-54.04	Office Supplies	

see notes for results

Merging DataFrames

```
pandas.merge(DataFrame1, DataFrame2, on="col_name", how="inner/outer/left/right")
```

DFs1

		id	type
0	SNP_FB_0411211	SNP	
1	SNP_FB_0412425	SNP	
2	SNP_FB_0942385	SNP	
3	CH01f09	SSR	
4	Hi05f12x	SSR	
5	SNP_FB_0942712	SNP	

DFs2

	chr	id
0	1	SNP_FB_0411211
1	15	SNP_FB_0412425
2	7	SNP_FB_0942385
3	9	CH01f09
4	1	SNP_FB_0428218

1. how = inner : non-matching entries are discarded;
2. how = left : ids are taken from the first DataFrame;
3. how = right : ids are taken from the second DataFrame;
4. how = outer : ids from both are retained.

```
inJ = pd.merge(DFs1, DFs2, on = "id", how = "inner")  
print(inJ)
```

Inner merge (only common in both)

	id	type	chr
0	SNP_FB_0411211	SNP	1
1	SNP_FB_0412425	SNP	15
2	SNP_FB_0942385	SNP	7
3	CH01f09	SSR	9

Right merge (IDS from DFs2)

	id	type	chr
0	SNP_FB_0411211	SNP	1
1	SNP_FB_0412425	SNP	15
2	SNP_FB_0942385	SNP	7
3	CH01f09	SSR	9
4	SNP_FB_0428218	NaN	1

```
leftJ = pd.merge(DFs1, DFs2, on = "id", how = "left")  
print(leftJ)
```

Left merge (IDS from DFs1)

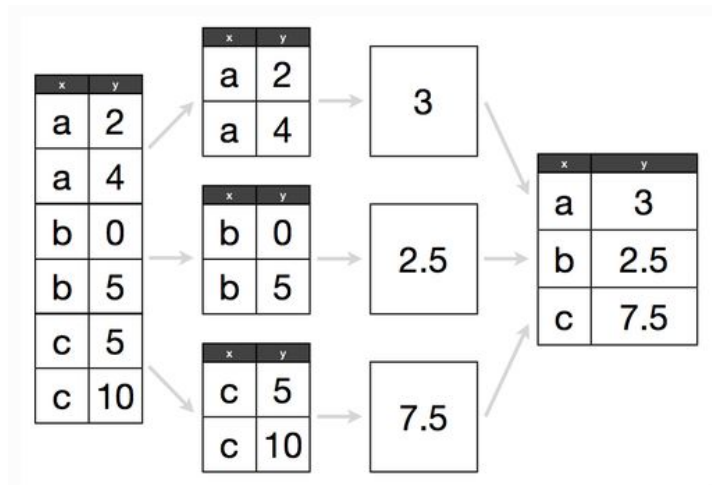
	id	type	chr
0	SNP_FB_0411211	SNP	1
1	SNP_FB_0412425	SNP	15
2	SNP_FB_0942385	SNP	7
3	CH01f09	SSR	9
4	Hi05f12x	SSR	NaN
5	SNP_FB_0942712	SNP	NaN

Outer merge (IDS from both)

	id	type	chr
0	SNP_FB_0411211	SNP	1
1	SNP_FB_0412425	SNP	15
2	SNP_FB_0942385	SNP	7
3	CH01f09	SSR	9
4	Hi05f12x	SSR	NaN
5	SNP_FB_0942712	SNP	NaN
6	SNP_FB_0428218	NaN	1

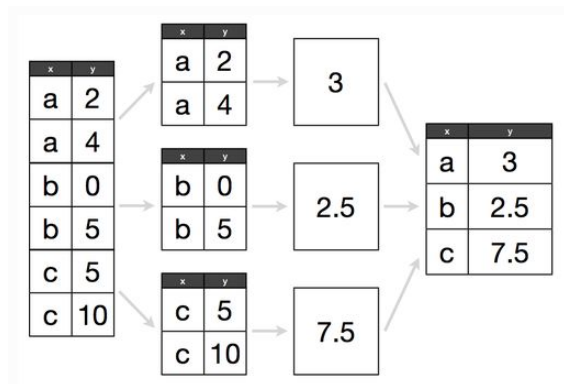
Grouping DataFrames

The split-apply-aggregate
paradigm



Grouping DataFrames

The split-apply-aggregate paradigm



```
import pandas as pd
```

```
test = {"x": ["a", "a", "b", "b", "c", "c"],  
        "y": [2, 4, 0, 5, 5, 10]}
```

```
DF = pd.DataFrame(test)
```

```
print(DF)
```

```
print("")
```

```
gDF = DF.groupby("x")
```

```
for i, g in gDF:
```

```
    print("Group: ", i)
```

```
    print(g)
```

```
    print(type(g))
```

```
aggDF = gDF.agg(lambda x: pd.DataFrame.mean(x))
```

```
print(aggDF)
```

	x	y
0	a	2
1	a	4
2	b	0
3	b	5
4	c	5
5	c	10

Group: a

	x	y
0	a	2
1	a	4

<class 'pandas.core.frame.DataFrame'>

Group: b

	x	y
2	b	0
3	b	5

<class 'pandas.core.frame.DataFrame'>

Group: c

	x	y
4	c	5
5	c	10

<class 'pandas.core.frame.DataFrame'>

y

x	y
a	3.0
b	2.5
c	7.5

Grouping DataFrames

Row ID	Sales	Profit	Product Category
1	261.5400	-213.25	Office Supplies
49	10123.0200	457.81	Office Supplies
50	244.5700	46.71	Office Supplies
80	4965.7595	1198.97	Technology
85	394.2700	30.94	Office Supplies

Group: Furniture
Group: Office Supplies
Group: Technology

Count elements per category:

Office Supplies	4610
Technology	2065
Furniture	1724

Name: Product Category, dtype: int64

Total values:

Product Category	Sales	Profit
Furniture	5178590.542	117433.03
Office Supplies	3752762.100	518021.42
Technology	5984248.182	886313.52

Mean values (sorted by profit):

Product Category	Sales	Profit
Furniture	3003.822820	68.116607
Office Supplies	814.048178	112.369072
Technology	2897.941008	429.207516

The most profitable is Technology

Questions:

How many Product categories?

Total sales and profits per category?

What is the most profitable category?

```
import pandas as pd
import matplotlib.pyplot as plt

orders = pd.read_csv("file_samples/sampledata_orders.csv", sep=";",
                     index_col=0, header=0)

SPC = orders[["Sales", "Profit", "Product Category"]]
print(SPC.head())

SPC.plot(kind="hist", bins=10)
plt.show()

print("")
grouped = SPC.groupby("Product Category")
for i, g in grouped:
    print("Group: ", i)

print("")
print("Count elements per category:") #get the series corresponding to the column
                                     #and apply the value_counts() method
print(orders["Product Category"].value_counts())
print("")
print("Total values:")
print(grouped.aggreate(pd.DataFrame.sum))

print("Mean values (sorted by profit):")
mv_sorted = grouped.aggreate(pd.DataFrame.mean).sort_values(by="Profit")
print(mv_sorted)
print("")
print("The most profitable is {}".format(mv_sorted.index[-1]))
```

<http://pandas.pydata.org/pandas-docs/stable/api.html>

pandas 0.20.3 documentation »

Table Of Contents

- What's New
- Installation
- Contributing to pandas
- Package overview
- 10 Minutes to pandas
- Tutorials
- Cookbook
- Intro to Data Structures
- Essential Basic Functionality
- Working with Text Data
- Options and Settings
- Indexing and Selecting Data
- Multindex / Advanced Indexing
- Computational tools
- Working with missing data
- Group By: split-apply-combine
- Merge, join, and concatenate
- Reshaping and Pivot Tables
- Time Series / Date functionality
- Time Deltas
- Categorical Data
- Visualization
- Styling
- IO Tools (Text, CSV, HDF5, ...)
- Remote Data Access
- Enhancing Performance
- Sparse data structures
- Frequently Asked Questions (FAQ)
- rpy2 / R interface
- pandas Ecosystem
- Comparison with R / R libraries
- Comparison with SQL
- Comparison with SAS
- API Reference
 - Input/Output
 - Pickling
 - pandas.read_pickle
 - Flat File
 - pandas.read_table

API Reference

This page gives an overview of all public pandas objects, functions and methods. In general, all classes and functions exposed in the top-level `pandas.*` namespace are regarded as public.

Further some of the subpackages are public, including `pandas.errors`, `pandas.plotting`, and `pandas.testing`. Certain functions in the `pandas.io` and `pandas.tseries` submodules are public as well (those mentioned in the documentation). Further, the `pandas.api.types` subpackage holds some public functions related to data types in pandas.

Warning: The `pandas.core`, `pandas.compat`, and `pandas.util` top-level modules are considered to be PRIVATE. Stability of functionality in those modules is not guaranteed.

Input/Output

Pickling

`read_pickle(path[, compression])` Load pickled pandas object (or any other pickled object) from the specified

Flat File

<code>read_table(filepath_or_buffer[, sep, ...])</code>	Read general delimited file into DataFrame
<code>read_csv(filepath_or_buffer[, sep, ...])</code>	Read CSV (comma-separated) file into DataFrame
<code>read_fwf(filepath_or_buffer[, colspecs, widths])</code>	Read a table of fixed-width formatted lines into DataFrame
<code>read_msgpack(path_or_buf[, encoding, iterator])</code>	Load msgpack pandas object from the specified

Clipboard

`read_clipboard([sep])` Read text from clipboard and pass to `read_table`.

Excel

First things first

We are going to need some libraries

```
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
```

In **Linux** you can install the libraries by typing in a terminal `sudo pip3 install matplotlib`, `sudo pip3 install pandas` and `sudo pip3 install numpy` (or `sudo python3.X -m pip install matplotlib`, `sudo python3.X -m pip install pandas` and `sudo python3.6 -m pip install numpy`), where X is your python version.

In **Windows** you can install the libraries by typing in the command prompt (to open it type `cmd` in the search) `pip3 install matplotlib`, `pip3 install pandas` and `pip3 install numpy`.

Exercises

1. The file [top_3000_words.txt](#) is a one-column file representing the top 3000 English words. Read the file and for each letter, count how many words start with that letter. Store this information in a dictionary. Create a pandas series from the dictionary and plot an histogram of all initials counting more than 100 words starting with them.

Show/Hide Solution

2. The file [filt_aligns.tsv](#) is a tab separated value file representing alignments of paired-end reads on some apple chromosomes. Paired end reads have the property of being X bases apart from each other as they have been sequenced from the two ends of some size-selected DNA molecules.



Each line of the file has the following information

`readID\tChrPE1\tAlignmentPosition1\tChrPE2\tAlignmentPosition2`. The two ends of the same pair have the same readID. Load the read pairs aligning on the same chromosome into two dictionaries. The first (`inserts`) having readID as keys and the insert size (i.e. the absolute value of `AlignmentPosition1 - AlignmentPosition2`) as value. The second dictionary (`chrs`) will have readID as key and chromosome ID as value. Example:

```
readID Chr11 31120 Chr11 31472
readID1 Chr7 12000 Chr11 11680
```

will result in:

```
inserts = {"readID" : 352, "readID1" : 320}
chrs = {"readID" : "Chr11", "readID1" : "Chr7"}
```